

Africa crisis

October 9, 2019

1 Problem setting

1.1 Africa Economic, Banking and Systemic Crisis :

Context : This dataset is a derivative of Reinhart et. al's Global Financial Stability dataset which can be found online at: <https://www.hbs.edu/behavioral-finance-and-financial-stability/data/Pages/global.aspx> The dataset will be valuable to those who seek to understand the dynamics of financial stability within the African context.

1.2 Content :

The dataset specifically focuses on the Banking, Debt, Financial, Inflation and Systemic Crises that occurred, from 1860 to 2014, in 13 African countries, including: Algeria, Angola, Central African Republic, Ivory Coast, Egypt, Kenya, Mauritius, Morocco, Nigeria, South Africa, Tunisia, Zambia and Zimbabwe. Acknowledgements Reinhart, C., Rogoff, K., Trebesch, C. and Reinhart, V. (2019) Global Crises Data by Country. [online] <https://www.hbs.edu/behavioral-finance-and-financial-stability/data>. Available at: <https://www.hbs.edu/behavioral-finance-and-financial-stability/data/Pages/global.aspx> [Accessed: 17 July 2019].

1.3 Inspiration :

Which factors are most associated with Systemic Crises in Africa ?

2 Dataset description :

2.1 Input variables :

- systemic_crisis "0" means that no systemic crisis occurred in the year and "1" means that a systemic crisis occurred in the year.
- exch_usd The exchange rate of the country vis-a-vis the USD
- domestic_debt_in_default "0" means that no sovereign domestic debt default occurred in the year and "1" means that a sovereign domestic debt default occurred in the year
- sovereign_external_debt_default "0" means that no sovereign external debt default occurred in the year and "1" means that a sovereign external debt default occurred in the year
- gdp_weighted_default The total debt in default vis-a-vis the GDP
- inflation_annual_cpi The annual CPI Inflation rate
- independence "0" means "no independence" and "1" means "independence"

- currency_crises "0" means that no currency crisis occurred in the year and "1" means that a currency crisis occurred in the year
- inflation_crises "0" means that no inflation crisis occurred in the year and "1" means that an inflation crisis occurred in the year ## Predict variable (desired target):

banking_crisis "no_crisis" means that no banking crisis occurred in the year and "crisis" means that a banking crisis occurred in the year

2.2 Introduce the Data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: # Import data and take a look
data= pd.read_csv("african_crises.csv")
data.head()
```

```
Out[2]:
```

	case	cc3	country	year	systemic_crisis	exch_usd	\
0	1	DZA	Algeria	1870	1	0.052264	
1	1	DZA	Algeria	1871	0	0.052798	
2	1	DZA	Algeria	1872	0	0.052274	
3	1	DZA	Algeria	1873	0	0.051680	
4	1	DZA	Algeria	1874	0	0.051308	

	domestic_debt_in_default	sovereign_external_debt_default	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	gdp_weighted_default	inflation_annual_cpi	independence	currency_crises	\
0	0.0	3.441456	0	0	
1	0.0	14.149140	0	0	
2	0.0	-3.718593	0	0	
3	0.0	11.203897	0	0	
4	0.0	-3.848561	0	0	

	inflation_crises	banking_crisis
0	0	crisis
1	0	no_crisis
2	0	no_crisis
3	0	no_crisis
4	0	no_crisis

```
In [3]: print(data.shape)
```

```
(1059, 14)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1059 entries, 0 to 1058
Data columns (total 14 columns):
case                1059 non-null int64
cc3                 1059 non-null object
country             1059 non-null object
year                1059 non-null int64
systemic_crisis     1059 non-null int64
exch_usd            1059 non-null float64
domestic_debt_in_default 1059 non-null int64
sovereign_external_debt_default 1059 non-null int64
gdp_weighted_default 1059 non-null float64
inflation_annual_cpi 1059 non-null float64
independence        1059 non-null int64
currency_crises     1059 non-null int64
inflation_crises    1059 non-null int64
banking_crisis      1059 non-null object
dtypes: float64(3), int64(8), object(3)
memory usage: 115.9+ KB
```

Our next step is to divide the train data into “attributes” and “labels”. X variable contains all the attributes/features and Y variable contains labels.

```
In [5]: # Assign X as a DataFrame of features and y as a Series of the outcome variable
X = data.drop('banking_crisis', 1)
y = data.banking_crisis
```

```
In [6]: X = X.drop('case', 1)
X = X.drop('cc3', 1)
```

```
In [7]: print(X.shape)
```

```
(1059, 11)
```

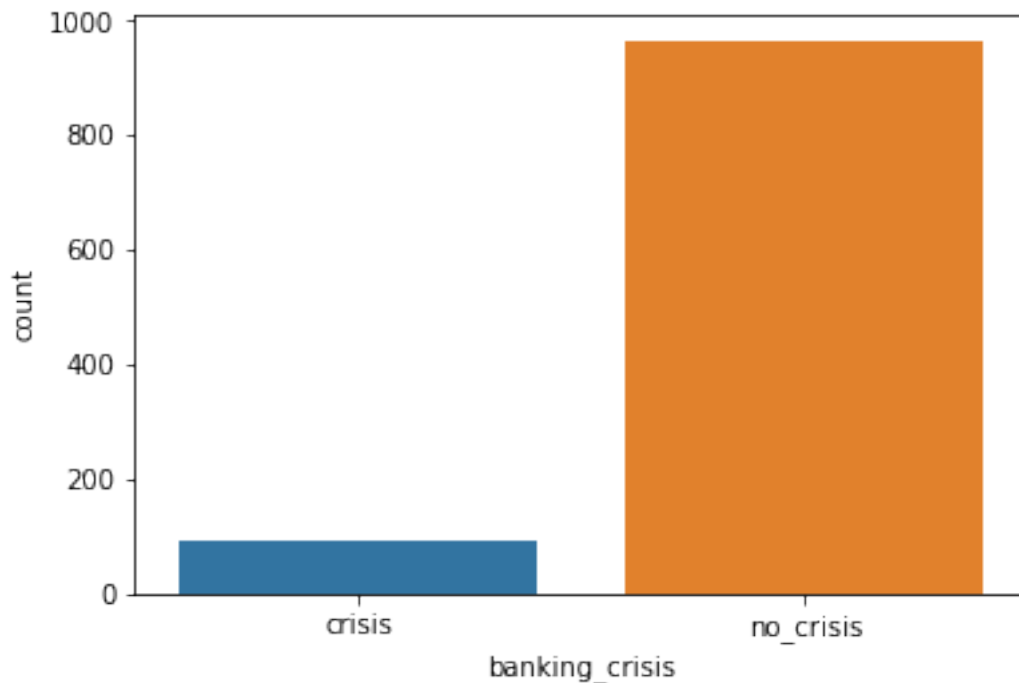
2.3 Exploring the data

Lets now explore the data with few visualizations.

```
In [8]: # Take a look at the outcome variable: 'income'
print(data['banking_crisis'].value_counts())
```

```
no_crisis    965
crisis        94
Name: banking_crisis, dtype: int64
```

```
In [9]: # Plotting the distribution of the labels in the bar plot
sns.countplot(x='banking_crisis', data=data);
```



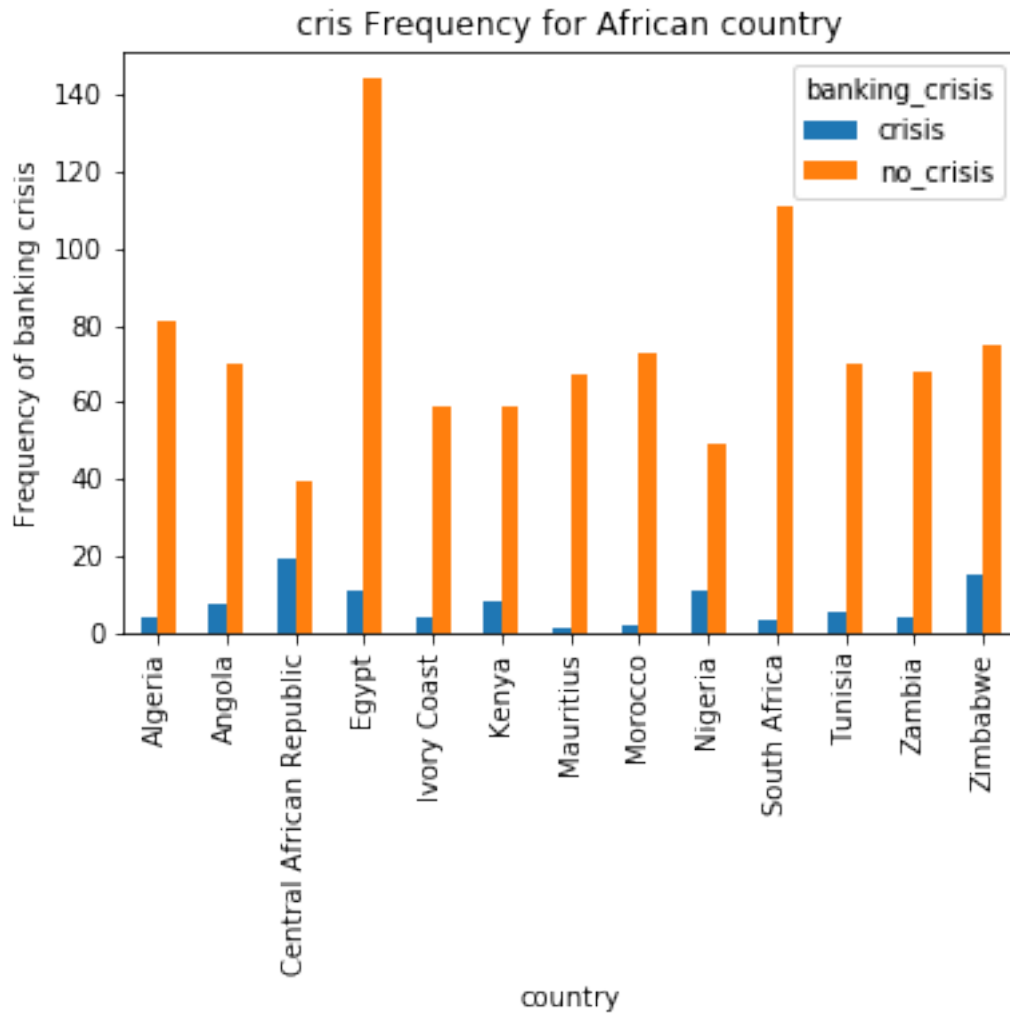
Our classes are imbalanced

```
In [10]: # Visualisations
```

```
# cris by country
```

```
pd.crosstab(data.country,y).plot(kind='bar')
plt.title('cris Frequency for African country')
plt.xlabel('country')
plt.ylabel('Frequency of bankelg crisis')
```

```
Out[10]: Text(0,0.5,'Frequency of banking crisis')
```



```
In [12]: fig, axs = plt.subplots(ncols=2, nrows=4, figsize=(20, 20))
plt.subplots_adjust(hspace=0.68)
fig.delaxes(axs[3][1])

# Employment type
wc_plot = sns.countplot(X['systemic_crisis'], ax=axs[0][0])
wc_plot.set_xticklabels(wc_plot.get_xticklabels(), rotation=40, ha="right")

# Gender
ge_plot = sns.countplot(X['domestic_debt_in_default'], ax=axs[0][1])
ge_plot.set_xticklabels(ge_plot.get_xticklabels(), rotation=72, ha="right")

# Education level
ed_plot = sns.countplot(X['sovereign_external_debt_default'], ax=axs[1][0])
ed_plot.set_xticklabels(ed_plot.get_xticklabels(), rotation=40, ha="right")
```

```

# Marital status
ms_plot = sns.countplot(X['independence'], ax=axes[1][1])
ms_plot.set_xticklabels(ms_plot.get_xticklabels(), rotation=40, ha="right")

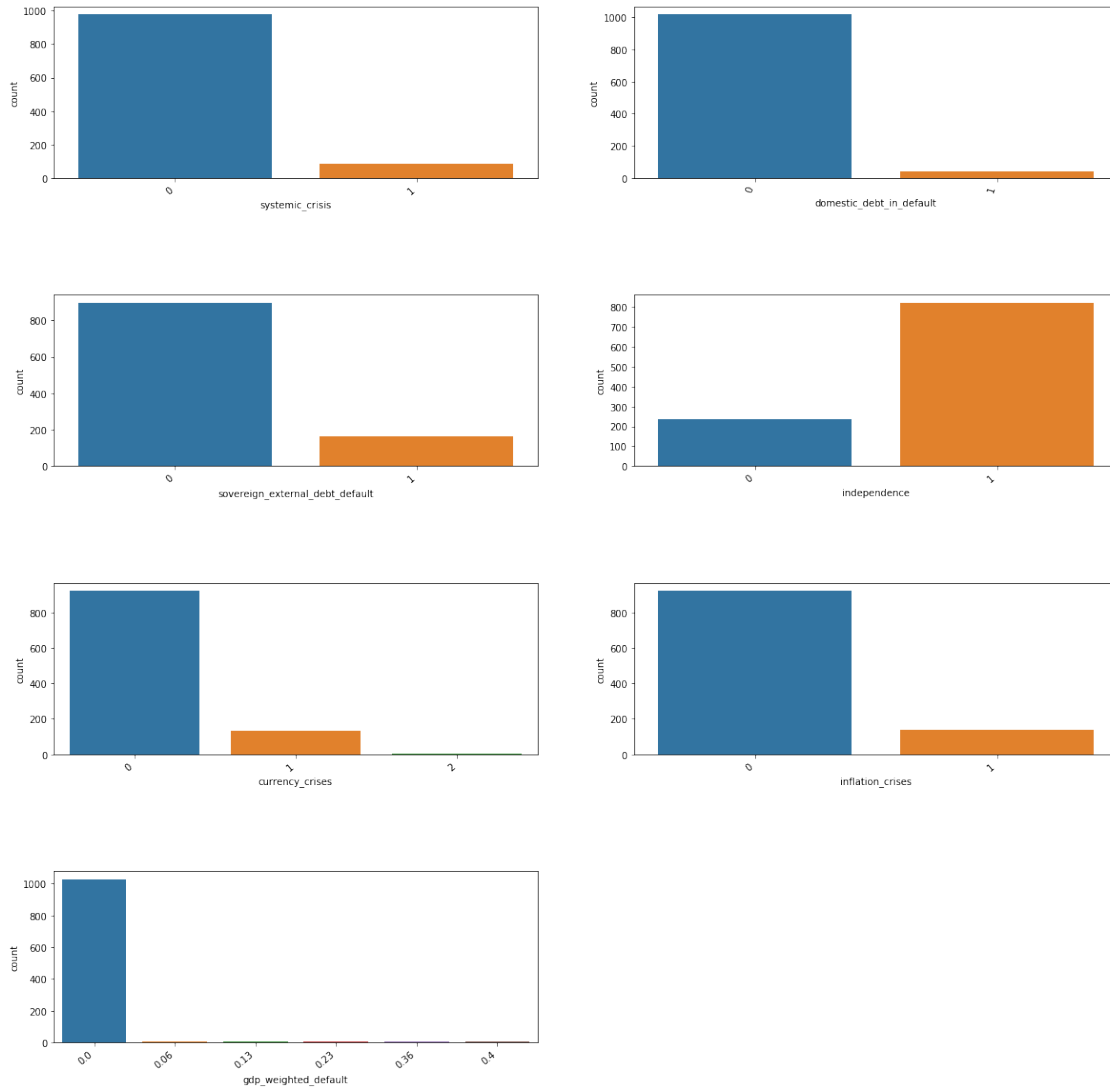
# Relationship
rel_plot = sns.countplot(X['currency_crises'], ax=axes[2][0])
rel_plot.set_xticklabels(rel_plot.get_xticklabels(), rotation=40, ha="right")

# Race
race_plot = sns.countplot(X['inflation_crises'], ax=axes[2][1])
race_plot.set_xticklabels(race_plot.get_xticklabels(), rotation=40, ha="right")

# Occupation
occ_plot = sns.countplot(X['gdp_weighted_default'], ax=axes[3][0])
occ_plot.set_xticklabels(occ_plot.get_xticklabels(), rotation=40, ha="right")

plt.show()

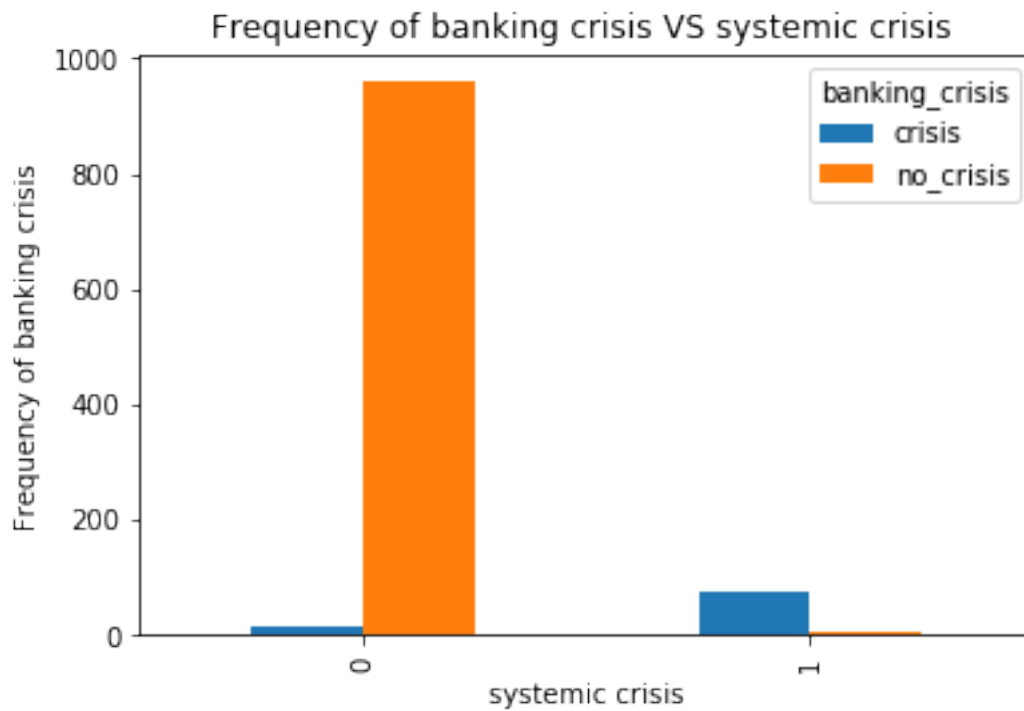
```



In [13]: *# visualisation for categorical variables*

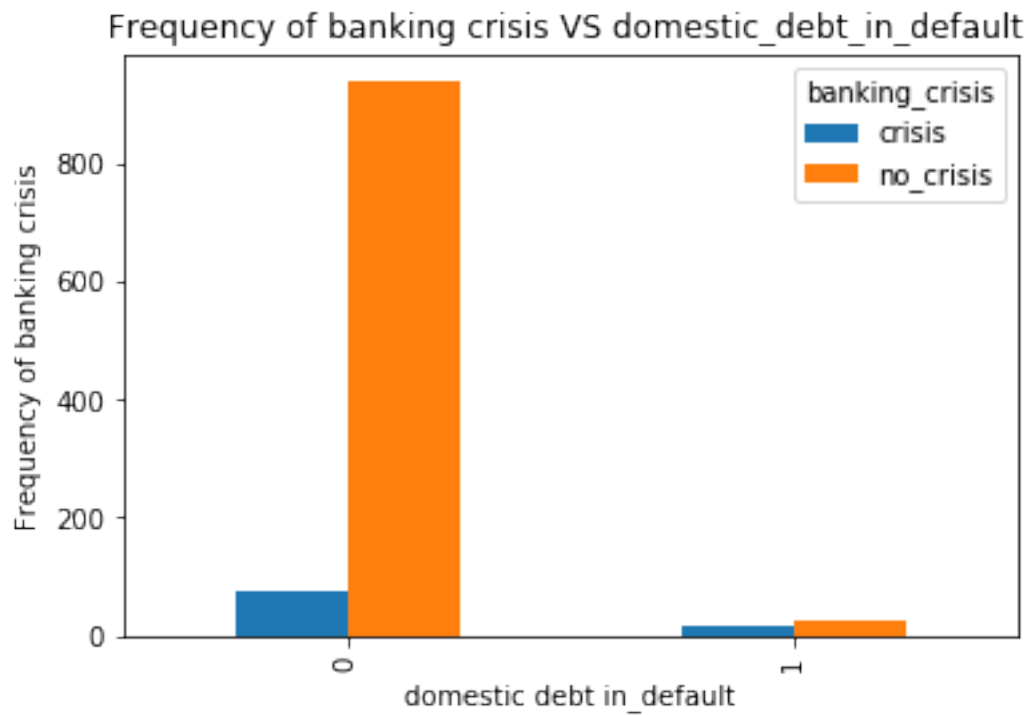
```
pd.crosstab(data.systemic_crisis,y).plot(kind='bar')
plt.title('Frequency of banking crisis VS systemic crisis ')
plt.xlabel('systemic crisis')
plt.ylabel('Frequency of banking crisis')
```

Out[13]: Text(0,0.5,'Frequency of banking crisis')



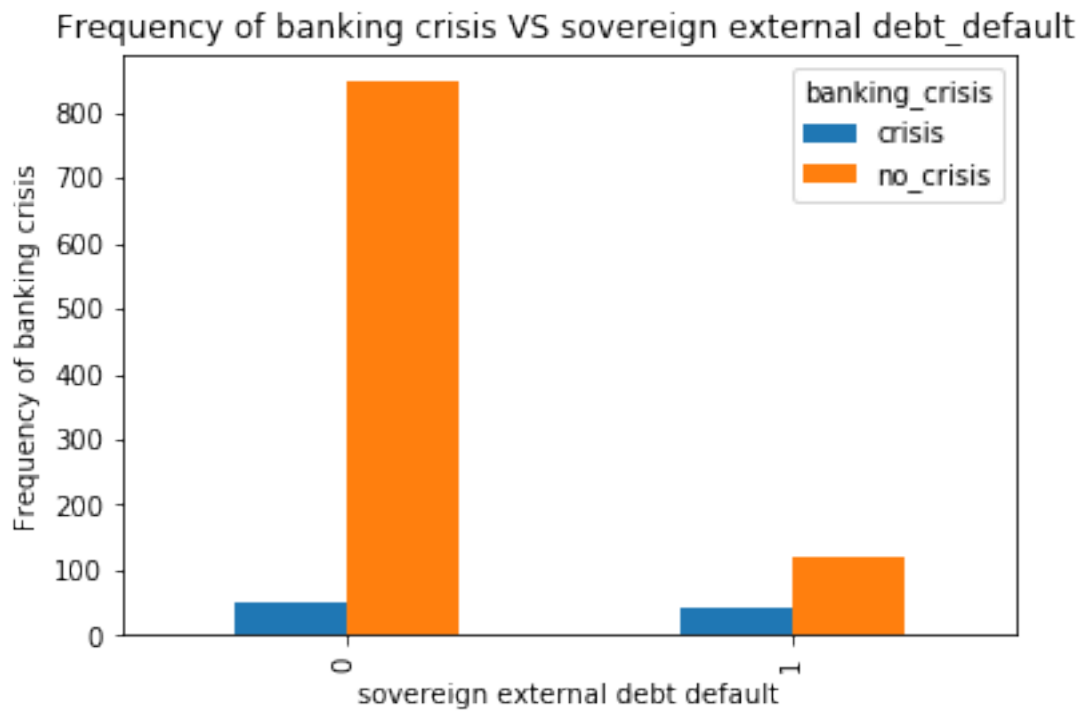
```
In [14]: pd.crosstab(data.domestic_debt_in_default,y).plot(kind='bar')
plt.title('Frequency of banking crisis VS domestic_debt_in_default ')
plt.xlabel('domestic debt in_default')
plt.ylabel('Frequency of banking crisis')
```

```
Out[14]: Text(0,0.5,'Frequency of banking crisis')
```

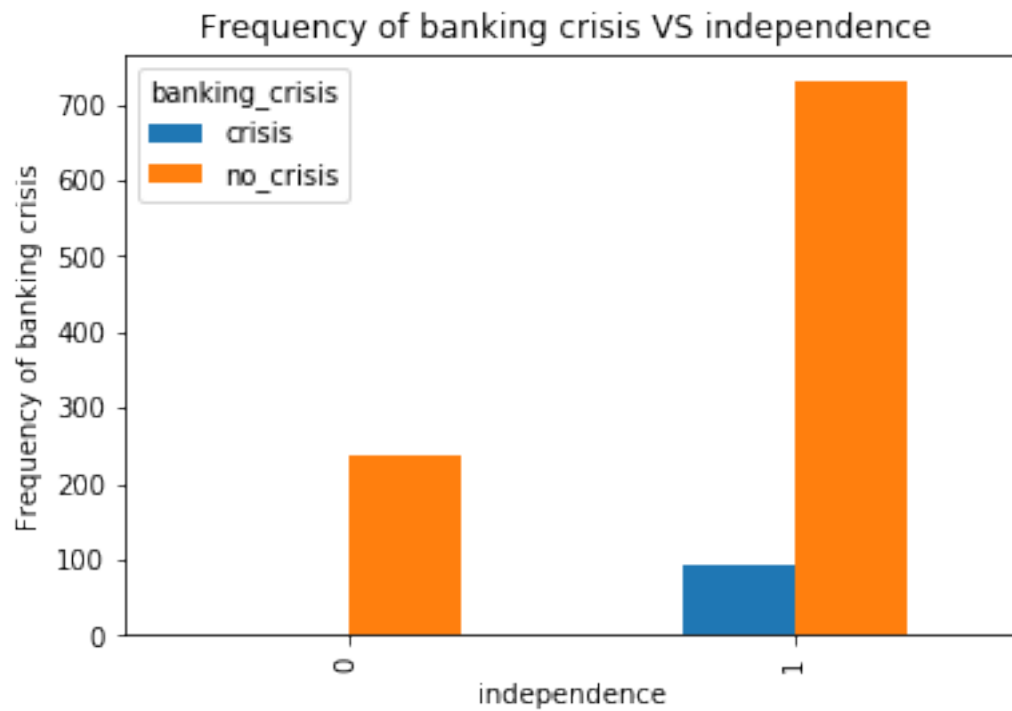
```
In [15]: pd.crosstab(data.sovereign_external_debt_default,y).plot(kind='bar')
plt.title('Frequency of banking crisis VS sovereign external debt_default ')
plt.xlabel('sovereign external debt default')
plt.ylabel('Frequency of banking crisis')
```

```
Out[15]: Text(0,0.5,'Frequency of banking crisis')
```



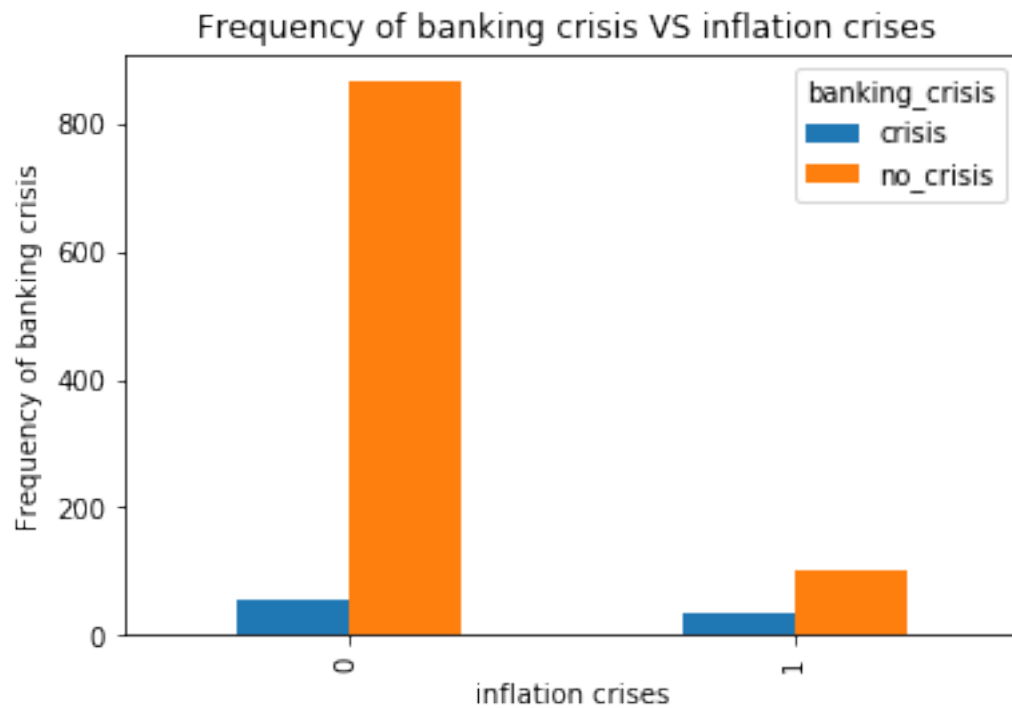
```
In [16]: pd.crosstab(data.independence,y).plot(kind='bar')
plt.title('Frequency of banking crisis VS independence ')
plt.xlabel('independence')
plt.ylabel('Frequency of banking crisis')
```

```
Out[16]: Text(0,0.5,'Frequency of banking crisis')
```



```
In [17]: pd.crosstab(data.inflation_crises,y).plot(kind='bar')
plt.title('Frequency of banking crisis VS inflation crises ')
plt.xlabel('inflation crises')
plt.ylabel('Frequency of banking crisis')
```

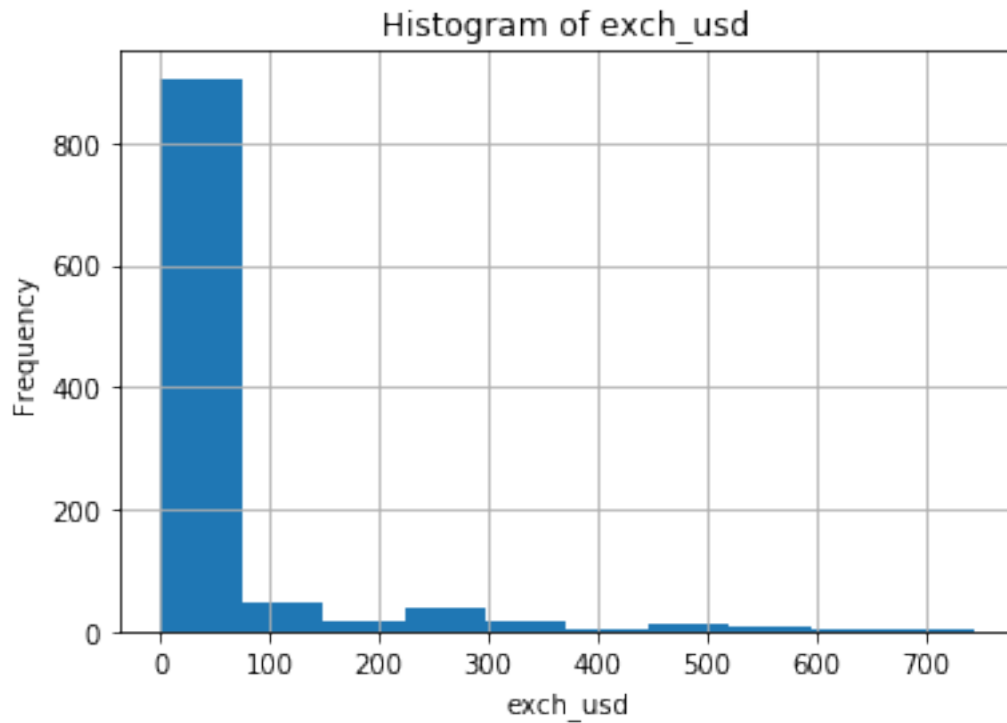
```
Out[17]: Text(0,0.5,'Frequency of banking crisis')
```



```
In [18]: # numerical variables
```

```
data.exch_usd.hist()  
plt.title('Histogram of exch_usd')  
plt.xlabel('exch_usd')  
plt.ylabel('Frequency')
```

```
Out[18]: Text(0,0.5,'Frequency')
```



2.3.1 Conclusion :

- Our result indicate that Systemic crises in Africa is significantly correleated with three factors :
 - systemic_crisis
 - exch_usd
 - independence

2.4 Data cleaning

2.4.1 Dealing with data types

Models can only handle numeric features Must convert categorical and ordinal features into numeric features

```
In [19]: # Decide which categorical variables you want to use in model
for col_name in X.columns:
    if X[col_name].dtypes == 'object':
        unique_cat = len(X[col_name].unique())
        print("Feature '{col_name}' has {unique_cat} unique categories".
              format(col_name=col_name, unique_cat=unique_cat))
```

Feature 'country' has 13 unique categories

```

In [20]: from sklearn.preprocessing import LabelEncoder
         le = LabelEncoder()
         for col_name in X.columns:
             if X[col_name].dtypes == 'object':
                 X[col_name]=le.fit_transform(X[col_name])

In [23]: # Our predict variable is categorical we have to convert it to binary .
         # Create dummy variables

         y= pd.get_dummies(y)
         y = y.drop(['no_crisis'], axis = 1)

In [24]: X.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1059 entries, 0 to 1058
Data columns (total 11 columns):
country                1059 non-null int64
year                  1059 non-null int64
systemic_crisis        1059 non-null int64
exch_usd              1059 non-null float64
domestic_debt_in_default 1059 non-null int64
sovereign_external_debt_default 1059 non-null int64
gdp_weighted_default   1059 non-null float64
inflation_annual_cpi   1059 non-null float64
independence          1059 non-null int64
currency_crises        1059 non-null int64
inflation_crises       1059 non-null int64
dtypes: float64(3), int64(8)
memory usage: 91.1 KB

In [25]: y['crisis'].value_counts()

Out[25]: 0    965
         1    94
         Name: crisis, dtype: int64

```

2.4.2 Handling missing data

An alternative solution is to use imputation - Replace missing value with another value - Strategies: mean, median, highest frequency value of given feature

```

In [26]: # How much of your data is missing?
         X.isnull().sum().sort_values(ascending=False).head()

         # Confirm All Missing Data is Handled

```

```
Out [26]: inflation_crises      0
          currency_crises      0
          independence          0
          inflation_annual_cpi  0
          gdp_weighted_default  0
          dtype: int64
```

```
In [27]: # Impute missing values using Imputer in sklearn.preprocessing
         #from sklearn.preprocessing import Imputer

         #imp = Imputer(missing_values='NaN', strategy='median', axis=0)
         #imp.fit(X)
         #X = pd.DataFrame(data=imp.transform(X) , columns=X.columns)
```

2.5 Model building

2.5.1 Build model using processed data

```
In [28]: # Importing objects from sklearn to help with the predictions.
         from sklearn.model_selection import train_test_split
         from sklearn.metrics import accuracy_score,
         precision_score, recall_score, confusion_matrix,
         precision_recall_curve
```

```
In [30]: # Use train_test_split in sklearn.cross_validation
         # to split data into train and test sets
         from sklearn.cross_validation import train_test_split

         X_train, X_test, y_train, y_test = train_test_split(X, y['crisis'],
         train_size=0.70, random_state=1)
```

C:\Users\awdii\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: "This module will be removed in 0.20.", DeprecationWarning)

```
In [31]: # Creating a confusion matrix.
         def CMatrix(CM, labels=['No crisis', 'crisis']):
             df = pd.DataFrame(data=CM, index=labels, columns=labels)
             df.index.name='TRUE'
             df.columns.name='PREDICTION'
             df.loc['TOTAL'] = df.sum()
             df['Total'] = df.sum(axis=1)
             return df
```

```
In [32]: # Preparing a pandas DataFrame to analyze models (evaluation metrics).
         metrics = pd.DataFrame(index=['accuracy', 'precision', 'recall'],
         columns=['NULL', 'LogisticReg', 'ClassTree', 'RandomForest'])
```

2.6 The Null Model : Always predict the most common category

```
In [33]: #####  
# The Null Model.  
y_pred_test = np.repeat(y_train.value_counts().idxmax(), y_test.size)  
metrics.loc['accuracy', 'NULL'] = accuracy_score(y_pred=y_pred_test, y_true=y_test)  
metrics.loc['precision', 'NULL'] = precision_score(y_pred=y_pred_test, y_true=y_test)  
metrics.loc['recall', 'NULL'] = recall_score(y_pred=y_pred_test, y_true=y_test)  
accuracy_score(y_pred=y_pred_test, y_true=y_test)
```

```
C:\Users\awdii\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1135: UndefinedMetricWarning: Precision is ill-defined for predicted class. The precision depends on whether the predicted values are considered as predicted or not.  
'precision', 'predicted', average, warn_for)
```

```
Out[33]: 0.9025157232704403
```

```
In [34]: CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)  
CMatrix(CM)
```

```
Out[34]: PREDICTION  No crisis  crisis  Total  
TRUE  
No crisis          287         0     287  
crisis              31         0      31  
TOTAL              318         0     318
```

2.7 A. Logistic Regression.

```
In [35]: # A. Logistic Regression.  
# 1- Import the estimator object (model).  
from sklearn.linear_model import LogisticRegression  
  
# 2- Create an instance of the estimator.  
logistic_regression = LogisticRegression()  
  
# 3- Use the training data to train the estimator.  
logistic_regression.fit(X_train, y_train)  
  
# 4- Evaluate the model.  
y_pred_test = logistic_regression.predict(X_test)  
metrics.loc['accuracy', 'LogisticReg'] = accuracy_score(y_pred=y_pred_test, y_true=y_test)  
metrics.loc['precision', 'LogisticReg'] = precision_score(y_pred=y_pred_test, y_true=y_test)  
metrics.loc['recall', 'LogisticReg'] = recall_score(y_pred=y_pred_test, y_true=y_test)  
  
# Confusion Matrix.  
CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)  
CMatrix(CM)
```

```
Out[35]: PREDICTION  No crisis  crisis  Total  
TRUE
```


No crisis	287	0	287
crisis	29	2	31
TOTAL	316	2	318

2.8 B. Classification Trees.

```
In [36]: # B. Classification Trees.
# 1- Import the estimator object (model).
from sklearn.tree import DecisionTreeClassifier

# 2- Create an instance of the estimator.
class_tree = DecisionTreeClassifier(min_samples_split=30,
                                    min_samples_leaf=10, random_state=10)

# 3- Use the training data to train the estimator.
class_tree.fit(X_train, y_train)

# 4- Evaluate the model.
y_pred_test = class_tree.predict(X_test)
metrics.loc['accuracy', 'ClassTree'] = accuracy_score(y_pred=y_pred_test, y_true=y_test)
metrics.loc['precision', 'ClassTree'] = precision_score(y_pred=y_pred_test, y_true=y_test)
metrics.loc['recall', 'ClassTree'] = recall_score(y_pred=y_pred_test, y_true=y_test)

# Confusion Matrix.
CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)
CMatrix(CM)
```

```
Out[36]: PREDICTION  No crisis  crisis  Total
TRUE
No crisis          284         3    287
crisis              5        26    31
TOTAL              289        29    318
```

2.9 C. Random Forest Classifier

```
In [37]: # C. Random Forest Classifier.
# 1- Import the estimator object (model).
from sklearn.ensemble import RandomForestClassifier

# 2- Create an instance of the estimator.
random_forest = RandomForestClassifier()

# 3- Use the training data to train the estimator.
random_forest.fit(X_train, y_train)

# 4- Evaluate the model.
y_pred_test = random_forest.predict(X_test)
```

```

metrics.loc['accuracy', 'RandomForest'] = accuracy_score(y_pred=y_pred_test, y_true=y_test)
metrics.loc['precision', 'RandomForest'] = precision_score(y_pred=y_pred_test, y_true=y_test)
metrics.loc['recall', 'RandomForest'] = recall_score(y_pred=y_pred_test, y_true=y_test)

```

```

# Confusion Matrix.

```

```

CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)
CMatrix(CM)

```

```

Out [37]: PREDICTION  No crisis  crisis  Total
TRUE
No crisis      285         2    287
crisis         4        27    31
TOTAL         289        29   318

```

```

In [38]: #####
# Comparing the models with percentages.
100*metrics

```

```

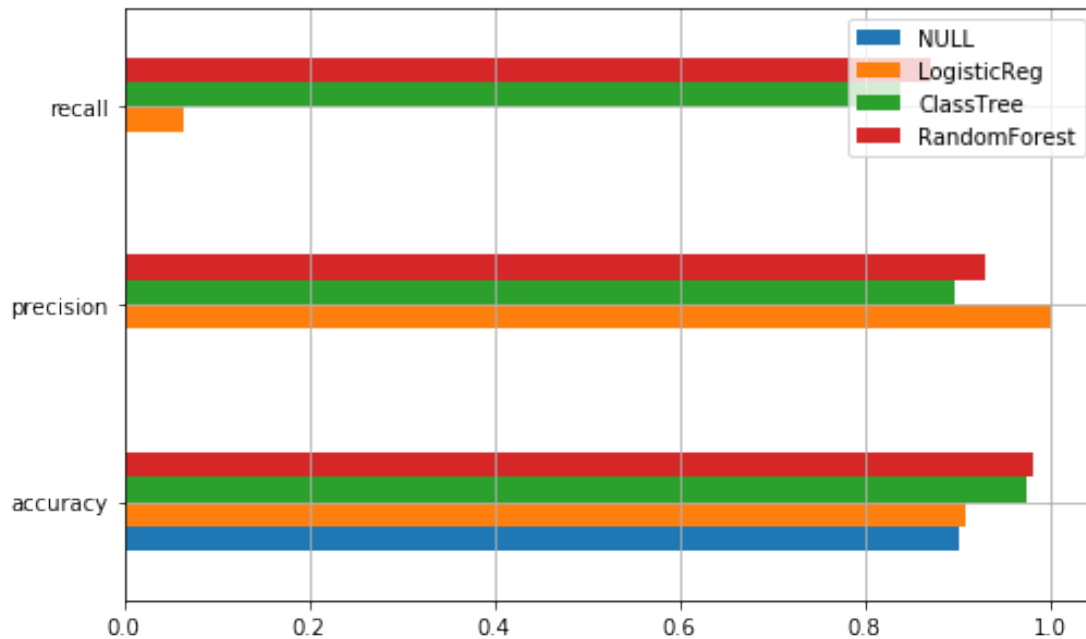
Out [38]:
          NULL LogisticReg ClassTree RandomForest
accuracy  90.2516    90.8805   97.4843    98.1132
precision      0        100   89.6552    93.1034
recall       0    6.45161   83.871    87.0968

```

```

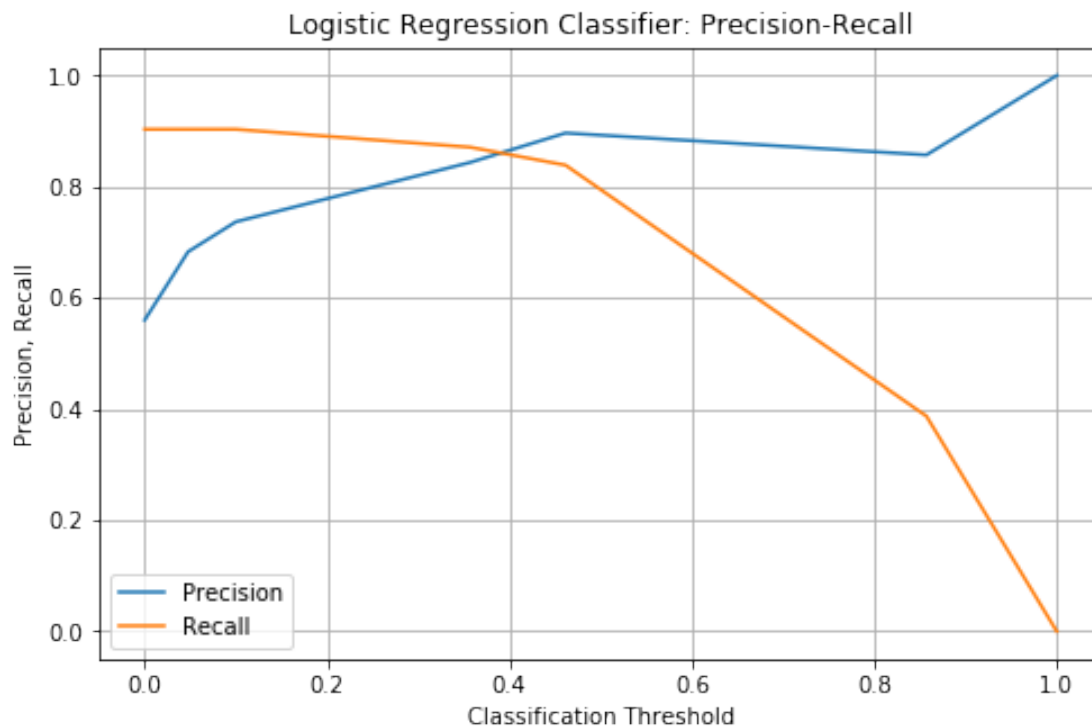
In [39]: # Comparing the models with a bar graph.
fig, ax = plt.subplots(figsize=(8,5))
metrics.plot(kind='barh', ax=ax)
ax.grid();

```



```
In [40]: # Adjusting the precision and recall values for the Descion Trees model.
precision_nb, recall_nb, thresholds_nb =
    precision_recall_curve(y_true=y_test,
        probas_pred=class_tree.predict_proba(X_test)[: ,1])
```

```
In [41]: # Creating a confusion matrix for modified Descion Trees model.
fig, ax = plt.subplots(figsize=(8,5))
ax.plot(thresholds_nb, precision_nb[1:], label='Precision')
ax.plot(thresholds_nb, recall_nb[1:], label='Recall')
ax.set_xlabel('Classification Threshold')
ax.set_ylabel('Precision, Recall')
ax.set_title('Logistic Regression Classifier: Precision-Recall')
ax.legend()
ax.grid();
```



2.10 Adjusting the threshold to 0.2.

```
In [42]: # Adjusting the threshold to 0.2.
y_pred_proba = class_tree.predict_proba(X_test)[: ,1]
y_pred_test = (y_pred_proba >= 0.2).astype('int')

# Confusion Matrix.
CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)
print('Recall: ', str(100*recall_score(y_pred=y_pred_test, y_true=y_test)) + '%')
```

```
print('Precision: ', str(100*precision_score(y_pred=y_pred_test, y_true=y_test)) + '%')
CMatrix(CM)
```

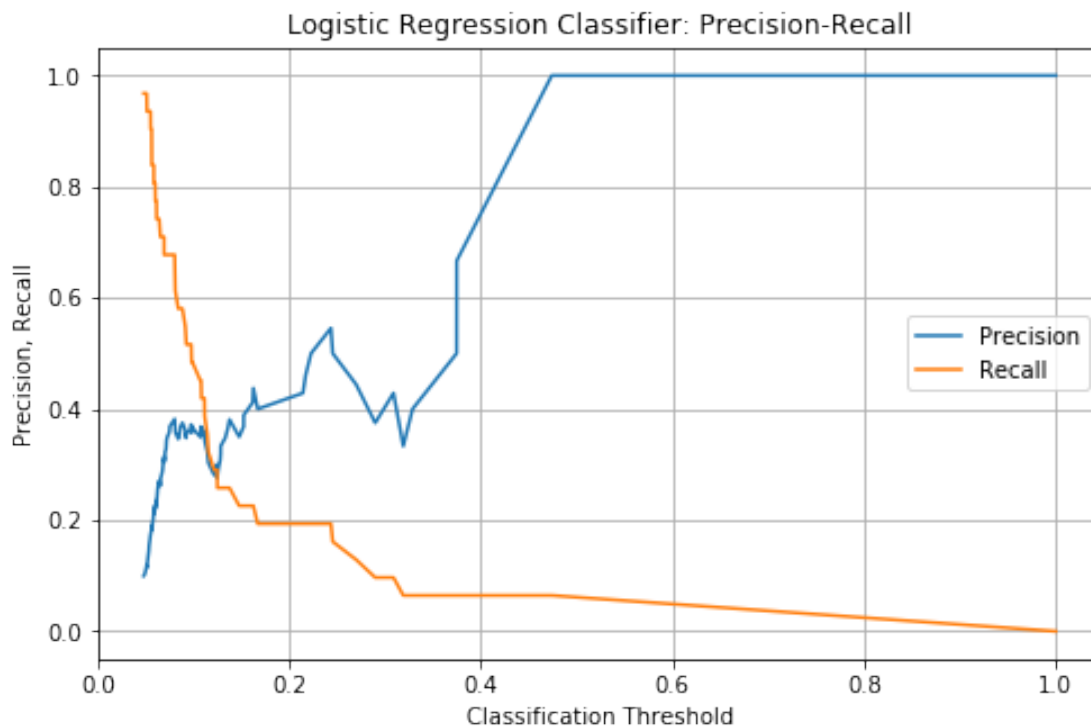
Recall: 90.32258064516128%

Precision: 73.68421052631578%

```
Out[42]: PREDICTION  No crisis  crisis  Total
TRUE
No crisis          277       10    287
crisis              3       28     31
TOTAL              280       38    318
```

```
In [43]: # Adjusting the precision and recall values for the logistic regression model.
precision_lr, recall_lr, thresholds_lr =
    precision_recall_curve(y_true=y_test,
        probas_pred=logistic_regression.predict_proba(X_test)[: ,1])
```

```
In [44]: # Creating a confusion matrix for modified Logistic Regression Classifier.
fig, ax = plt.subplots(figsize=(8,5))
ax.plot(thresholds_lr, precision_lr[1:], label='Precision')
ax.plot(thresholds_lr, recall_lr[1:], label='Recall')
ax.set_xlabel('Classification Threshold')
ax.set_ylabel('Precision, Recall')
ax.set_title('Logistic Regression Classifier: Precision-Recall')
ax.legend()
ax.grid();
```



2.11 Adjusting the threshold to 0.07

```
In [45]: # Adjusting the threshold to 0.07
```

```
y_pred_proba = logistic_regression.predict_proba(X_test)[:,-1]  
y_pred_test = (y_pred_proba >= 0.07).astype('int')
```

```
# Confusion Matrix.
```

```
CM = confusion_matrix(y_pred=y_pred_test, y_true=y_test)
```

```
print('Recall: ', str(100*recall_score(y_pred=y_pred_test, y_true=y_test)) + '%')
```

```
print('Precision: ', str(100*precision_score(y_pred=y_pred_test, y_true=y_test)) + '%')
```

```
CMatrix(CM)
```

```
Recall: 67.74193548387096%
```

```
Precision: 30.434782608695656%
```

```
Out[45]: PREDICTION  No crisis  crisis  Total  
TRUE  
No crisis          239        48    287  
crisis              10        21     31  
TOTAL              249        69    318
```